

Interpreting regression output

Contents

1	Elements of the typical regression output summary	2
2	Transformations	5
3	Interpreting coefficients of a regression with...	6
3.1	... Log transformations	6
3.2	... Interacted predictors	6
4	Post-estimation model diagnostics	7

*Disclaimer: sections and lines in brown correspond to content which is **very much** ‘under construction’.*

1 Elements of the typical regression output summary

```
> mod1 = lm(dist ~ speed, data = cars)
> summary(mod1)
```

Residuals:

Min	1Q	Median	3Q	Max
-29.069	-9.525	-2.272	9.215	43.201

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	42.9800	2.1750	19.761	< 2e-16 ***
speed.c	3.9324	0.4155	9.464	1.49e-12 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 15.38 on 48 degrees of freedom
Multiple R-squared: 0.6511, Adjusted R-squared: 0.6438
F-statistic: 89.57 on 1 and 48 DF, p-value: 1.49e-12

Figure 1: Summary of the results of the OLS estimation of a simple linear regression model, in R

Regression is a tool for estimating average differences across groups. We estimate the difference in average observed outcomes.

Residuals $\{r_i\}_i$ Difference between the observed response values y_i and those predicted \hat{y}_i .

Errors $\varepsilon_i = y_i - X_i\beta$, residuals are estimates of the errors: $r_i = y_i - X_i\hat{\beta} = y_i - \hat{y}_i$

Plot them to look at their distribution: is it centered around 0, is it normal...?

Coefficients For each coefficient, an estimate and the level of uncertainty for that estimate.

- **Slope estimate $\hat{\beta}_j$**

- With multiple predictors, the interpretation of any given coefficient is, in part, contingent on the other variables in the model. → Interpret each coefficient “with all the other predictors held constant”.

Furthermore, if predictors are correlated, it is important to note the following distinction: $\hat{\beta}_j$ measures not the *total* change in y expected from increasing x_j , but the *additional* change from increasing x_j , when the effects of all other variables are already accounted for.

- **Interpretation as comparison:** slope coefficients should be interpreted as *comparisons between* units that differ in one predictor:

“ $\hat{\beta}_j$ = how y differs, on average, when comparing two groups of units that differ by 1 in the predictor x_j .” Interpretation as changes *within* units, i.e. a “counterfactual” interpretation “ $\hat{\beta}_j$ = the expected change in y caused by adding 1 to x_j ” requires justification other than the data, e.g., in causal inference studies.

- x_j continuous: $\hat{\beta}_j$ = average change in y for a 1-unit change in x_j , holding other x ’s constant

- x_j categorical: e.g. binary: $\hat{\beta}_j$ = average difference in y between the category for which $x_j = 0$ and the category for which $x_j = 1$

$$\hat{\beta} = \frac{\sum_i (x_i - \bar{x})(y_i - \bar{y})}{\sum_i (x_i - \bar{x})^2}$$

If there is a single regressor, $\hat{\beta} = \frac{\text{cov}[x, y]}{\text{var}[x]}$

- (estimated) Standard Error $\widehat{\text{SE}}(\hat{\beta}_j) = \frac{\hat{\sigma}_{\hat{\beta}_j}}{\sqrt{n}}$

= an estimate of the standard deviation of $\hat{\beta}_j$'s sampling distribution.

The SE gives us a sense of our uncertainty about $\hat{\beta}_j$: the expected difference in $\hat{\beta}_j$ if we were to run the model again and again. A lower SE *relative to the coefficient* means more certainty.

SEs are used in computing confidence intervals and in the t -statistic for hypothesis testing.

- **t-statistic** $t_{\hat{\beta}} = \frac{\hat{\beta} - h_0}{\widehat{\text{SE}}(\hat{\beta})}$

The realization of the t-statistic $t_{\beta} = \frac{\beta - h_0}{\text{SE}(\beta)}$, for our data, and the null hypothesis $H_0: \beta_j = 0$.

t_{β} can be used in a two-sided¹ t-test of H_0 , as it follows a Student's t -distribution under H_0 : $t_{\beta} \underset{h_0}{\sim} \mathcal{T}_{n-2}$.

If its realization $t_{\hat{\beta}}$ falls in the tails of that distribution, that would mean it is very unlikely given H_0 , therefore we can reject H_0 . We will examine that with the two-sided p-value.

- **p-value** = $\Pr[\text{observing a } T > |t_{\hat{\beta}}| \text{ under } H_0]$

I.e., the probability of observing data as extreme as that actually observed, assuming H_0 .²

p-value small (< 0.05) $\iff t_{\hat{\beta}}$ falls in the tail of the Student's \mathcal{T} -distribution
 \implies observing our $t_{\hat{\beta}}$ is highly unlikely under H_0
 \implies reject H_0
 \implies there is a relationship between y and x , $\hat{\beta}$ is “significant”.

Residual Standard Error or Standard Error of the Regression (SER)

Summary of the scale of the residuals. It is the average distance by which an observed value falls from the regression line, i.e., the accuracy to which the model can predict y . Interpretations:

- y can deviate from the true regression line by 15.38 ft, on average
- the model can predict y to an accuracy of about 15.38 points.
- about 68% of y will be within 15.38 of the predicted value.
- SER^2 = the variance “unexplained” by the model: the amount of variation remaining in y after we remove the variation due to x .

R^2 or coefficient of determination

How much better the sample data is fit by the sample regression line $y = \alpha + \beta X$ than by the sample mean

¹ By default, statistical packages carry out a two-sided test and therefore report the two-sided p-value; however we could also use the t-statistic to carry out a one-sided test.

² \triangle The p-value is often misinterpreted to be the probability that H_0 is true, when it is the probability of observing data as extreme or more extreme than that actually observed, assuming H_0 . $p\text{-value} = \Pr(\text{obs} \mid \text{hyp}) \neq \Pr(\text{hyp} \mid \text{obs})$

line, $y = \bar{Y}$. It is one way of measuring the goodness-of-fit. See Figure 2.

$$R^2 = 1 - \frac{\text{sum of squared residuals (SSR)}}{\text{total sum of squares (TSS)}} = 1 - \frac{\sum_i (y_i - \hat{y}_i)^2}{\sum_i (y_i - \bar{y})^2} \in [0, 1]$$

- In most linear models, $TSS = ESS + RSS$ (where ESS is the explained sum of squares $\sum_i (\hat{y}_i - \bar{y})^2$), therefore $R^2 = \frac{ESS}{TSS}$ is the proportion of the sample variance in y that is explained by X .
- ⚠ R^2 mechanically increases as more predictors are included in the regression. → Use the **adjusted R^2** which adjusts for the number of predictors.

F statistic for an F-test³ of overall significance.

H0: all coefficients equal 0 (no relationship between y and X). H1: $\beta_j \neq 0$, for at least one j . We test the full model against a model with no regressors.

F's p-value small \iff at least some of the parameters are nonzero and the regression equation does have some validity in fitting the data (the X 's are not purely random w.r.t. y)

$$F = \frac{\text{mean regression sum of squares (MSR)}}{\text{mean error sum of squares (MSE)}} = \frac{\frac{ESS}{k}}{\frac{SSR}{n-k-1}} \in [0, +\infty[$$

³ In general, an F-statistic is a ratio of two quantities, F-test tests the H0 that the quantities are roughly equal, i.e. $F \simeq 1$. Reject H0 if F high ($\gg 1$). How large F really needs to be depends on the number of data points and predictors: if large sample, F-stat slightly above 1 is sufficient to reject H0; if small sample, need a large F-stat.

2 Transformations

Inverse hyperbolic sine (IHS) transformation For outcomes that have a thick right tail, the standard solution is to take a log transformation⁴; it brings extreme values closer to the middle, so they don't have such a large effect on the results. However, when the outcome also has many zero-valued observations (e.g., wealth), natural log transformations don't work well as $\ln(0)$ is undefined.

Instead one can use the inverse hyperbolic sine (IHS or arcsinh) transformation:

$$\log(y_i + (y_i^2 + 1)^{\frac{1}{2}})$$

It approximates the natural logarithm (except for very small values of y , it is $\approx \log(2y_i) = \log(2) + \log(y_i)$), and so it can be interpreted in exactly the same way as a standard log-transformed dependent variable, but is defined at zero, thus allows retaining zero-valued observations.

⁴ Another solution is to run quantile regressions and analyze each part of the distribution separately.

3 Interpreting coefficients of a regression with...

3.1 ... Log transformations

Regression model	<i>Given a change in x, what change do we expect in y?</i>
level-level (linear) $y = \beta_0 + \beta_1 x + e$	If x increases by 1 unit, y increases by β_1 units.
log-level (log-linear) $\ln(y) = \beta_0 + \beta_1 x + e$	<p>If x increases by 1 unit, $\ln(y)$ increases by β_1 units, i.e. y increases by a factor e^{β_1}.</p> <ul style="list-style-type: none"> • if small $\hat{\beta}$: can approximate: y increases by $(100 \times \beta_1)\%$ • if large $\hat{\beta}$: approximation invalid. y increases by $\times e^{\beta_1}$
level-log $y = \beta_0 + \beta_1 \ln(x) + e$	If $\ln(x)$ increases by 1 unit, y increases by β_1 units, i.e. if x increases by 1%, y increases by $\frac{\beta_1}{100}$ units.
log-log $\ln(y) = \beta_0 + \beta_1 \ln(x) + e$	If x increases by 1%, y increases by $\beta_1\%$. (β_1 is an <i>elasticity</i> .)

Why can we interpret natural log changes as percentage changes?

The log function is approximately linear around 1, i.e., it is reasonable to do a first order Taylor approximation of $\ln(x)$ around $x = 1$:


$$\begin{aligned}
 f(x) &\simeq f'(1)(x - 1) + f(1) \\
 \ln(x) &\simeq \frac{1}{1}(x - 1) + 0 \\
 \ln(x) &\simeq x - 1
 \end{aligned}$$

Then a *small* difference in logs of y can be approximately interpreted as a percentage change in y :

$$\ln(y_2) - \ln(y_1) = \ln\left(\frac{y_2}{y_1}\right) \simeq \frac{y_2}{y_1} - 1 = \frac{y_2 - y_1}{y_1}$$

Therefore, in a log-linear regression $\ln(y) = \beta_0 + \beta_1 x + e$:

- if $\hat{\beta}_1$ small, one can say “A 1-unit increase in x corresponds to a $(100 \times \hat{\beta}_1)\%$ increase in y ”;
- if $\hat{\beta}_1$ large, one should stick to saying “A 1-unit increase in x corresponds to a $e^{\hat{\beta}_1}$ factor increase in y ”.



3.2 ... Interacted predictors

Adding an interaction term allows the slope to vary across subgroups, and changes the interpretation of all coefficients along the way. Examples:

- $\text{kid_score} = \beta_0 + \beta_1 \text{mom_hs} + \beta_2 \text{mom_iq} + \beta_3 \text{mom_hs} : \text{mom_iq}$
 β_3 represents the difference in the slope for *mom_iq*, comparing children with mothers who did and did not complete high school.
- $y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_3 x_1 x_2$ The effect of x_1 is $\beta_1 + \beta_3 x_2$: it is different for each value of x_2 .

4 Post-estimation model diagnostics

4 plots of the residuals

1. **“Normal QQ” plot** = *Are the residuals normally distributed?*

Distribution of the residuals relative to a normal distribution. Is the curve of residuals a straight line?

2. **“Scale-Location” plot** = *Are the residuals homoskedastic?*

Plot of $\sqrt{|r_i|}$ against fitted values \hat{y}_i . Is the vertical spread of points uniform along x (i.e., do residuals have a uniform variance across the range of predicted values)?

3. **“Residuals vs Leverage” plot** = *Influential Observations*

residuals against leverages.

- First we derive the hat (or “projection”) matrix from our linear model.
- We compute each observation’s “Cook’s distance” in $[0:1]$: the amount by which \hat{y}_i would change if the observation was omitted. This gives us the leverage of that observation.
- We draw a half-normal plot, sorting the observations by their leverage.

⇒ This shows whether some observations have a great influence on the regression (i.e., are leverage points). If we can’t see the Cook’s distance lines, it means that all points are well inside them: there are no leverage points.

4. **“Residuals vs Fitted” plot** = *Is there a non-linear pattern?*

Residuals against fitted values. Are the residuals equally spread around a horizontal line? If not, the model is missing a non-linear relationship between predictor variables and the outcome variable — which is therefore left out in the residuals.




Figure 2: Representation of the terms of the coefficient of determination $R^2 = 1 - \frac{SSR}{TSS}$. The red areas represent the squared residuals w.r.t. to the average value \bar{y} , the blue areas represent the squared residuals w.r.t. the linear regression. The better the linear regression fits the data in comparison to the simple average, the higher the R^2 .

Source: Orzetto - <https://commons.wikimedia.org/w/index.php?curid=11398293>